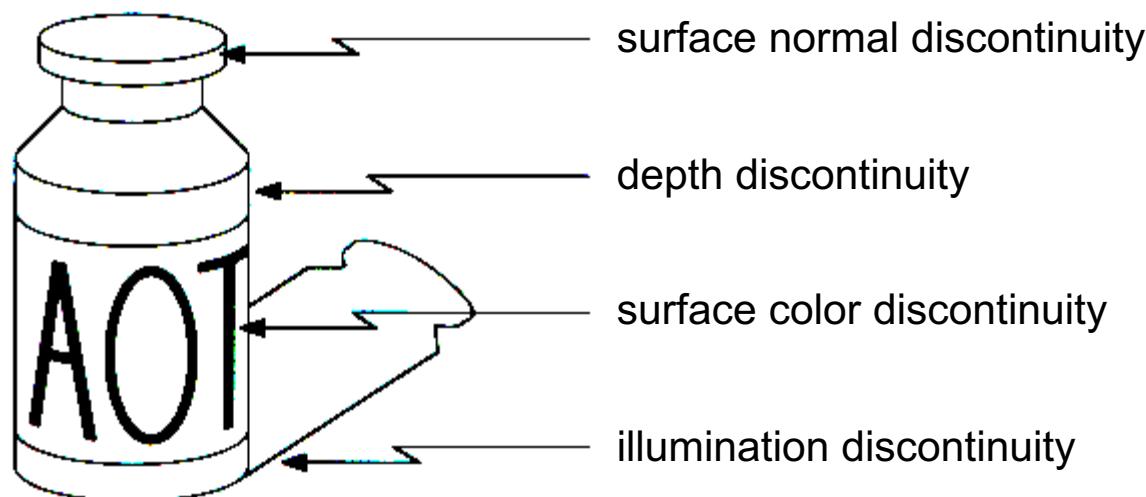

Edge Detection

CS 543 / ECE 549 – Saurabh Gupta
Spring 2020, UIUC

<http://saurabhg.web.illinois.edu/teaching/ece549/sp2020/>

Edge detection

- **Goal:** Identify sudden changes (discontinuities) in an image
- Intuitively, edges carry most of the semantic and shape information from the image



Edge detection

- **Ideal:** artist's line drawing

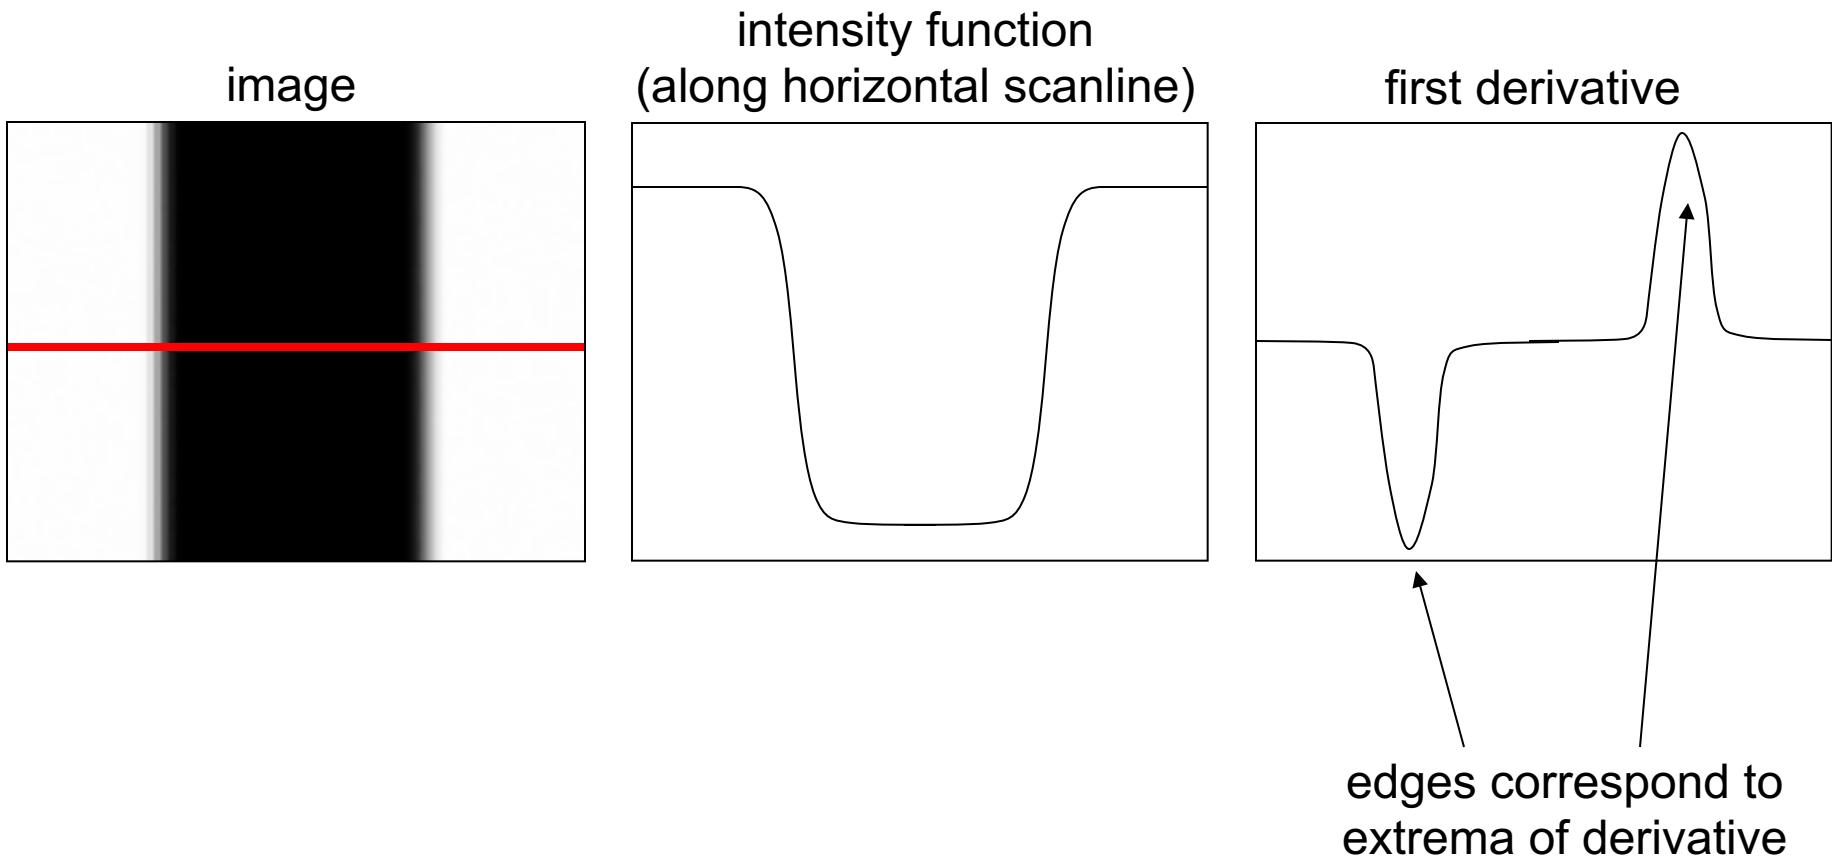


- **Reality:**



Edge detection

- An edge is a place of rapid change in the image intensity function



Derivatives with convolution

For 2D function $f(x,y)$, the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

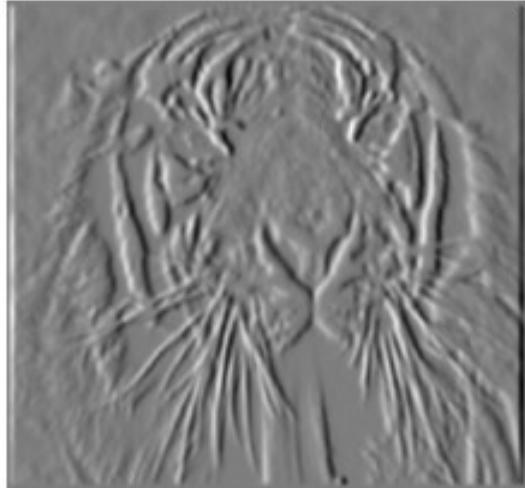
To implement the above as convolution, what would be the associated filter?

Partial derivatives of an image



$$\frac{\partial f(x, y)}{\partial x}$$

-1	1
----	---

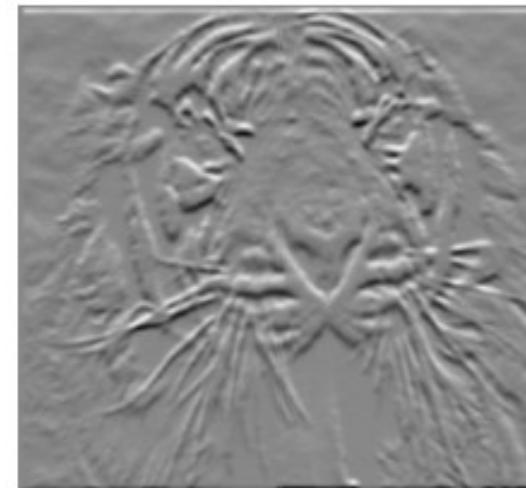


$$\frac{\partial f(x, y)}{\partial y}$$

-1	1
1	-1

 or

1	-1
-1	1



Which shows changes with respect to x?

Finite difference filters

Other approximations of derivative filters exist:

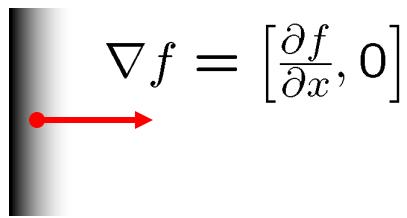
Prewitt: $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$

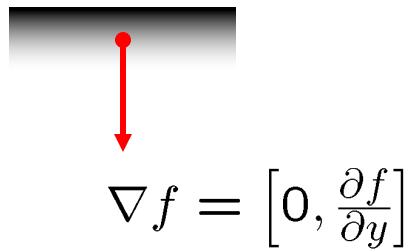
Sobel: $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$

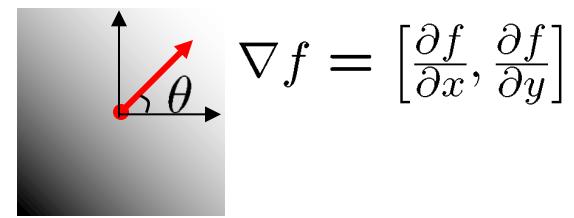
Roberts: $M_x = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$

Image gradient

The gradient of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$


$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$


$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$


$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient points in the direction of most rapid increase in intensity

- How does this direction relate to the direction of the edge?

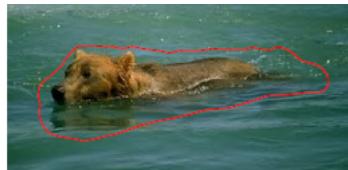
The gradient direction is given by $\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$

The edge strength is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2}$$

Application: Gradient-domain image editing

- Goal: solve for pixel values in the target region to match gradients of the source region while keeping background pixels the same



sources/destinations



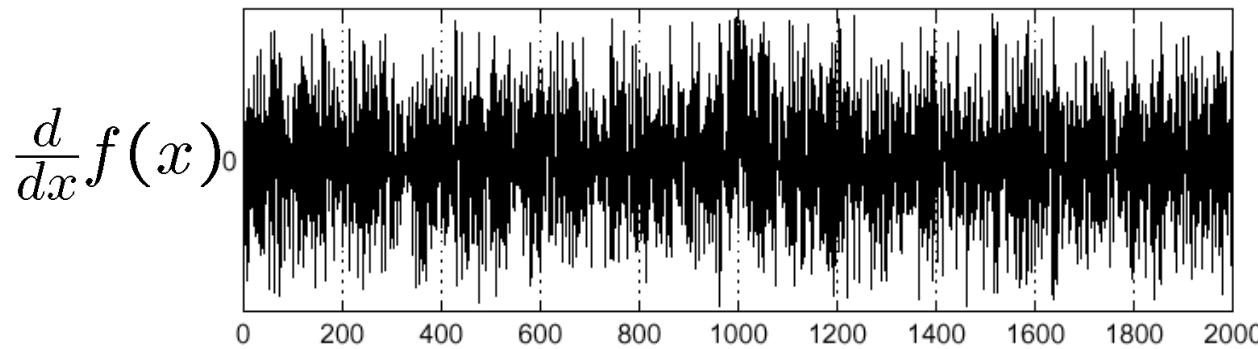
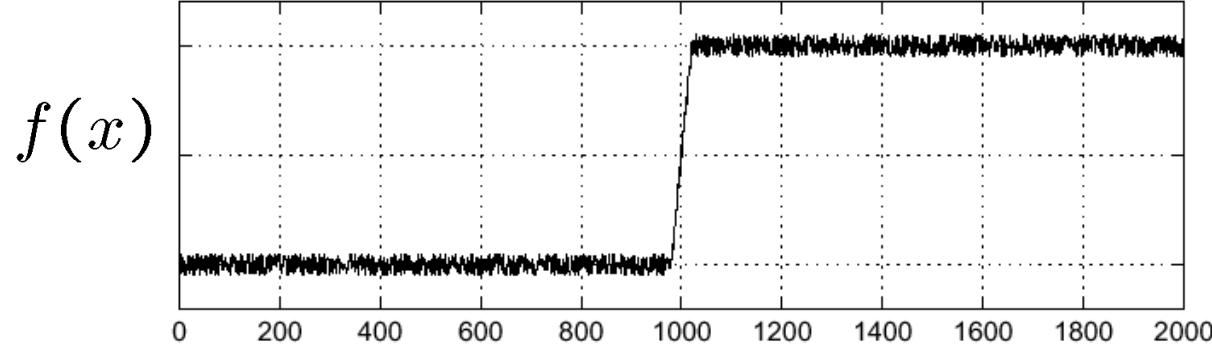
cloning



seamless cloning

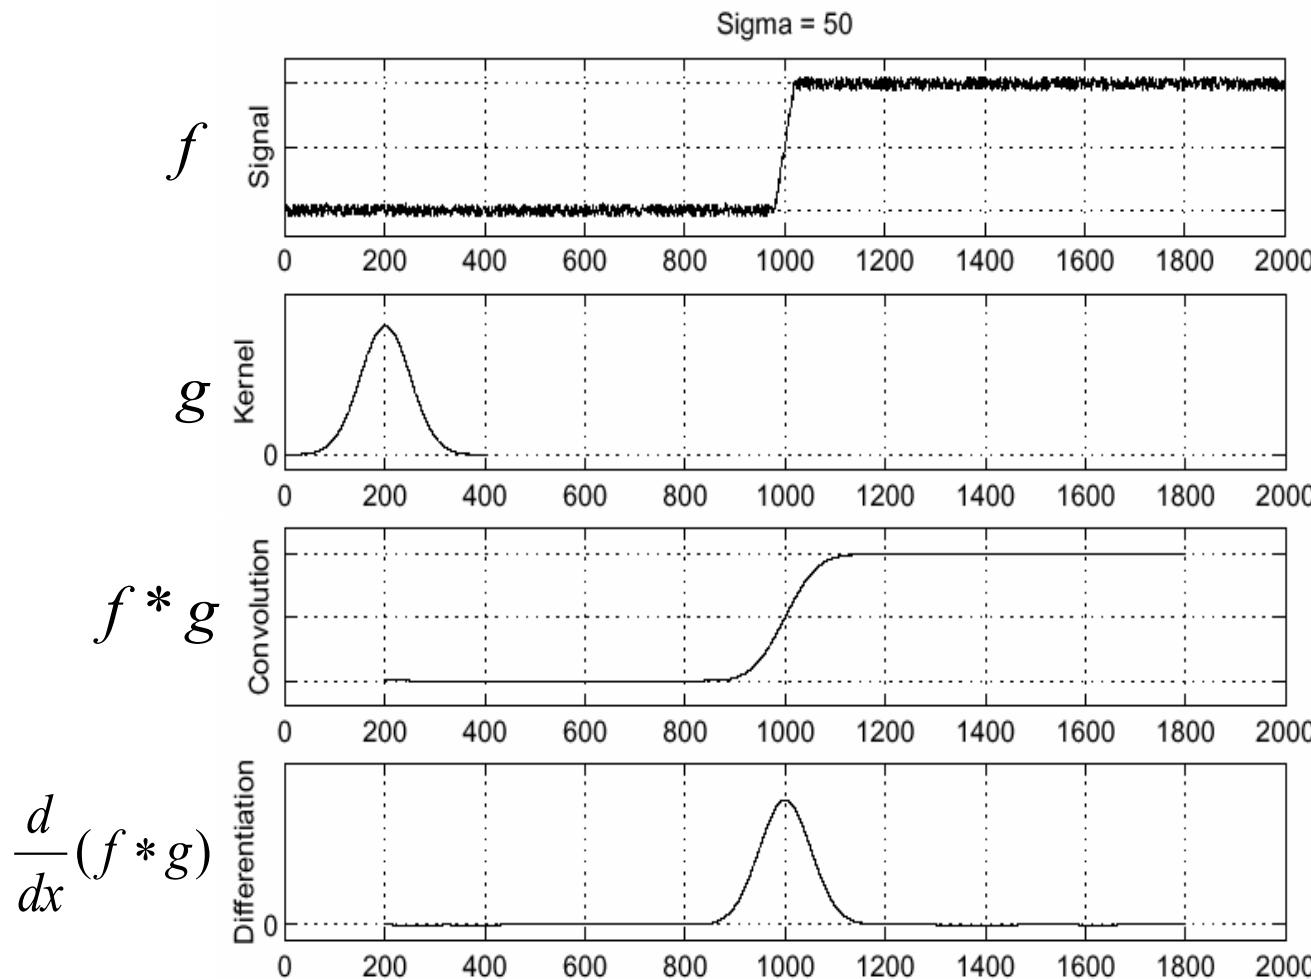
Effects of noise

Consider a single row or column of the image



Where is the edge?

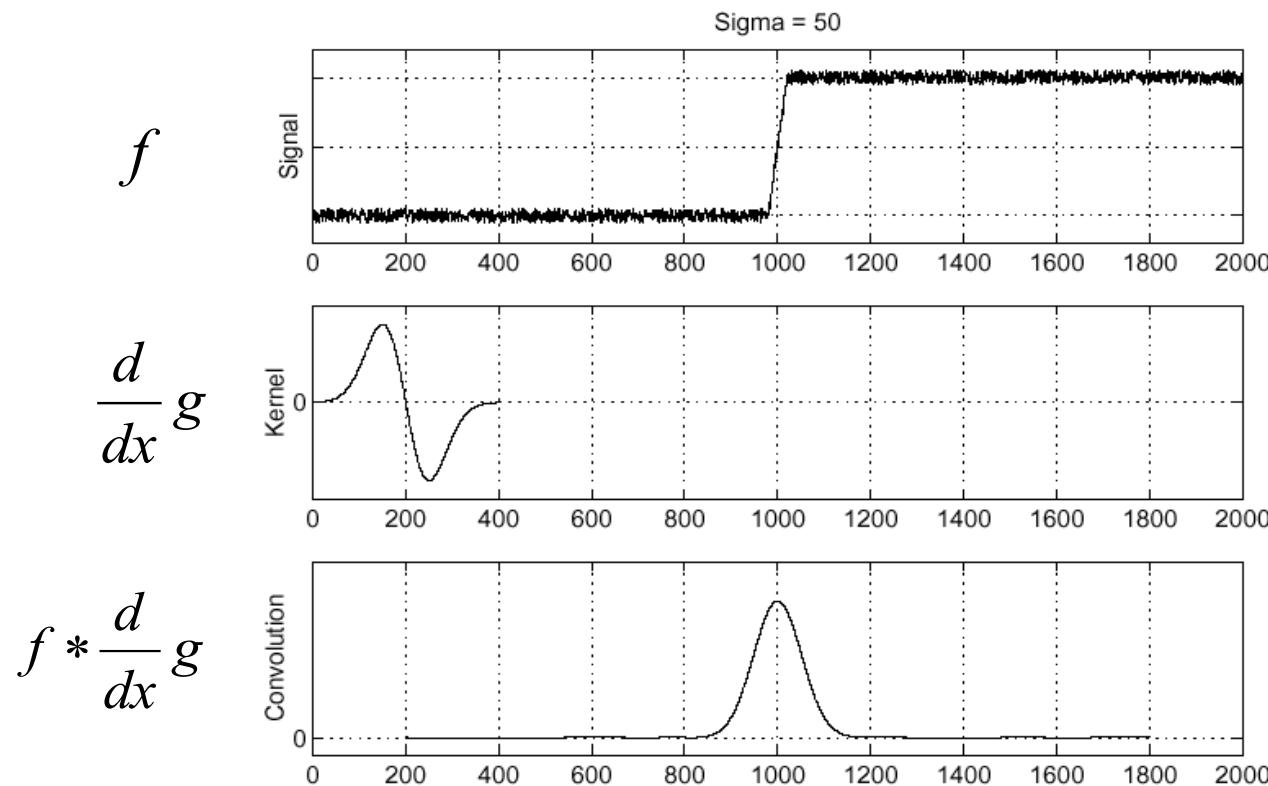
Solution: smooth first



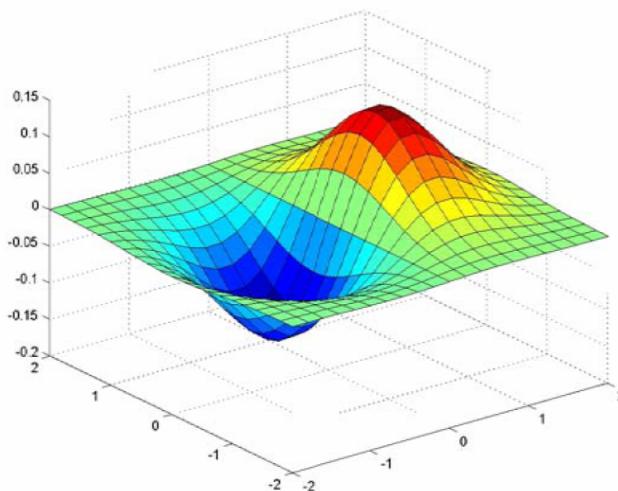
- To find edges, look for peaks in $\frac{d}{dx}(f * g)$

Derivative theorem of convolution

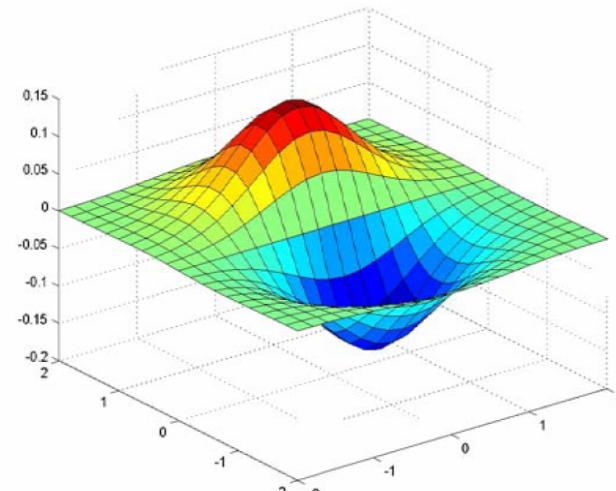
- Differentiation is convolution, and convolution is associative: $\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$
- This saves us one operation:



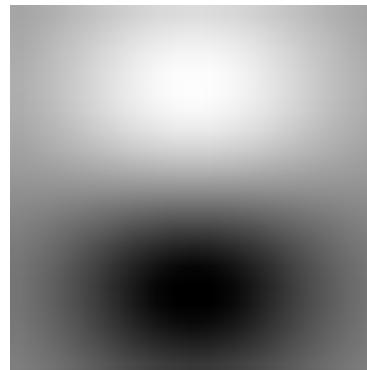
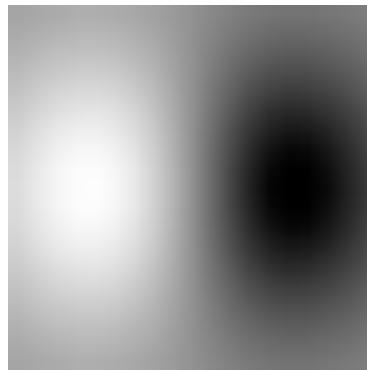
Derivative of Gaussian filters



x-direction

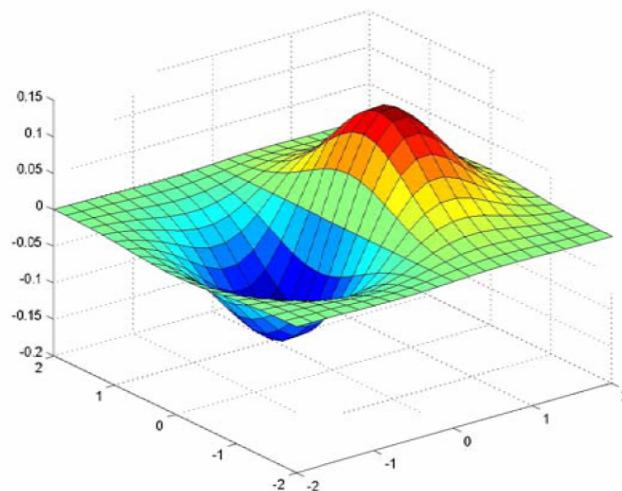


y-direction

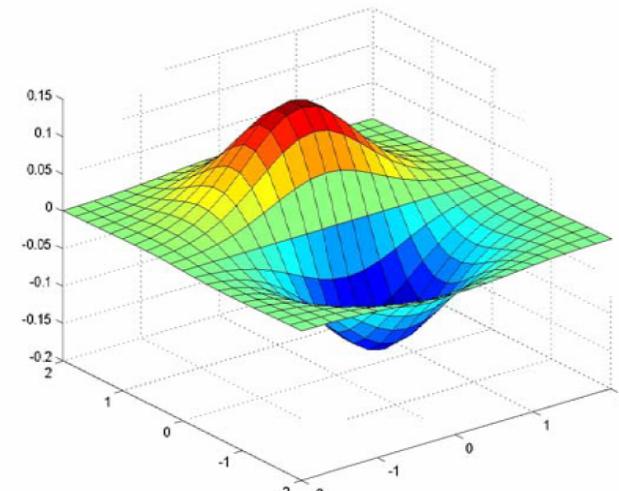


Which one finds horizontal/vertical edges?

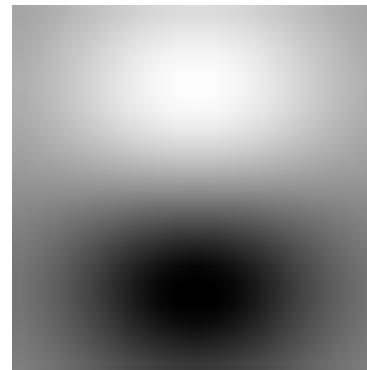
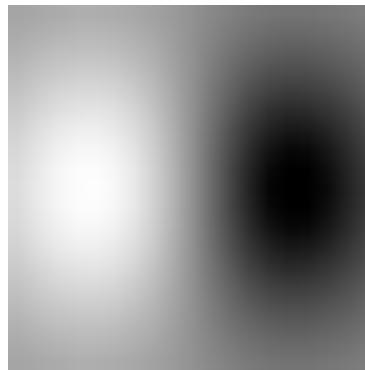
Derivative of Gaussian filters



x-direction



y-direction



Are these filters separable?

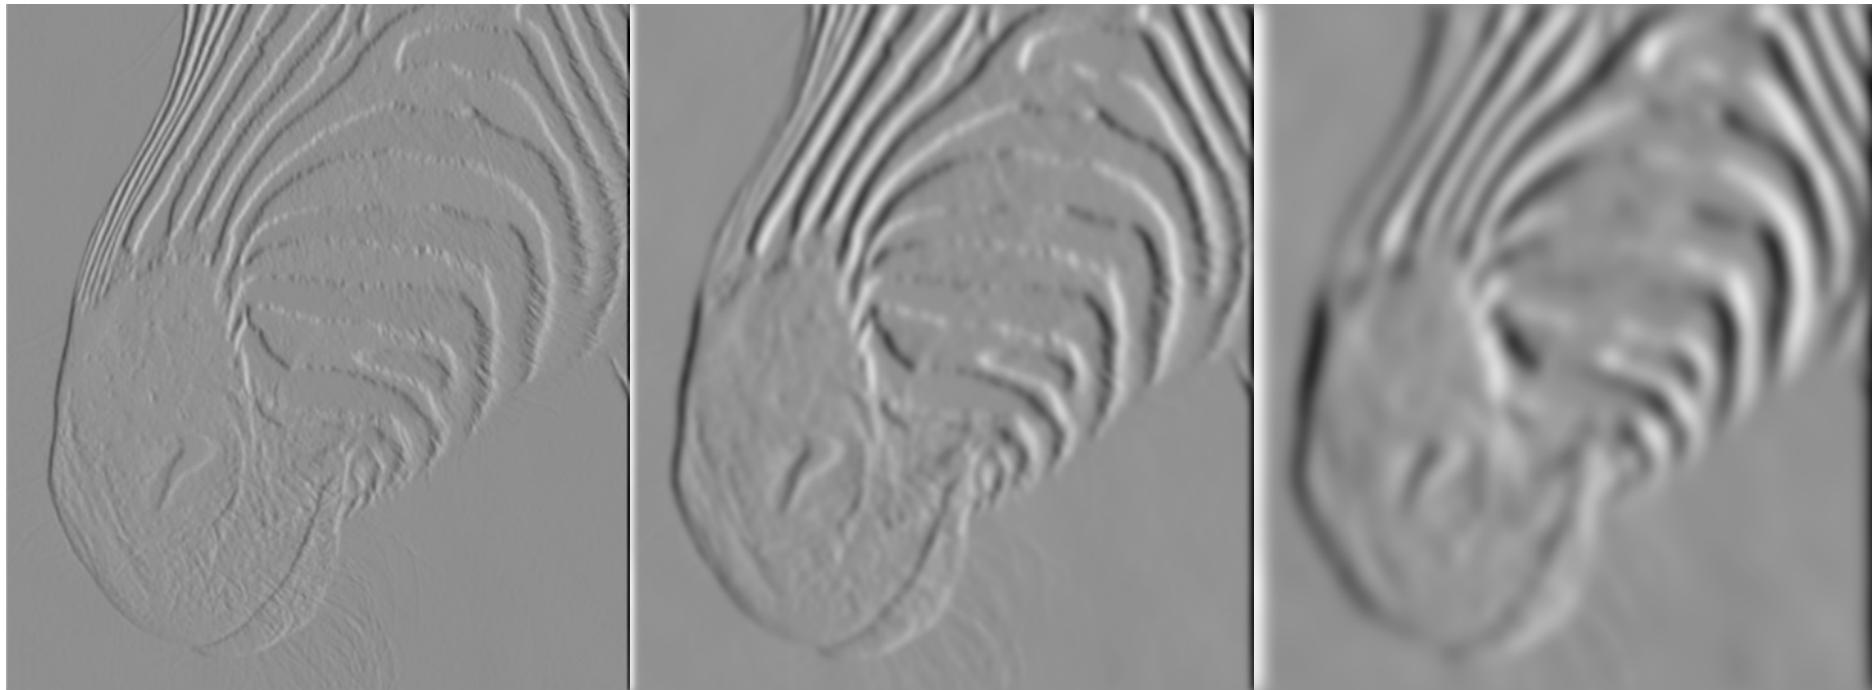
Recall: Separability of the Gaussian filter

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y

In this case, the two functions are the (identical) 1D Gaussian

Scale of Gaussian derivative filter



1 pixel

3 pixels

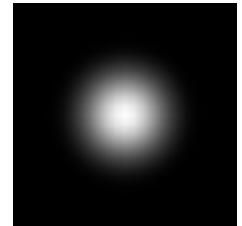
7 pixels

Smoothed derivative removes noise, but blurs edge. Also finds edges at different “scales”

Review: Smoothing vs. derivative filters

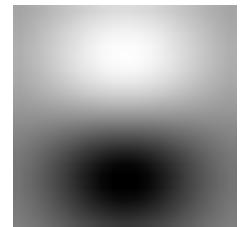
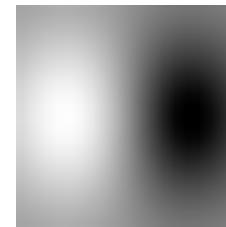
Smoothing filters

- Gaussian: remove “high-frequency” components; “low-pass” filter
- Can the values of a smoothing filter be negative?
- What should the values sum to?
 - **One**: constant regions are not affected by the filter



Derivative filters

- Derivatives of Gaussian
- Can the values of a derivative filter be negative?
- What should the values sum to?
 - **Zero**: no response in constant regions



Building an edge detector



original image



final output

Building an edge detector



norm of the gradient

Building an edge detector



$\theta = \text{atan2}(-gy, gx)$

orientation of the gradient

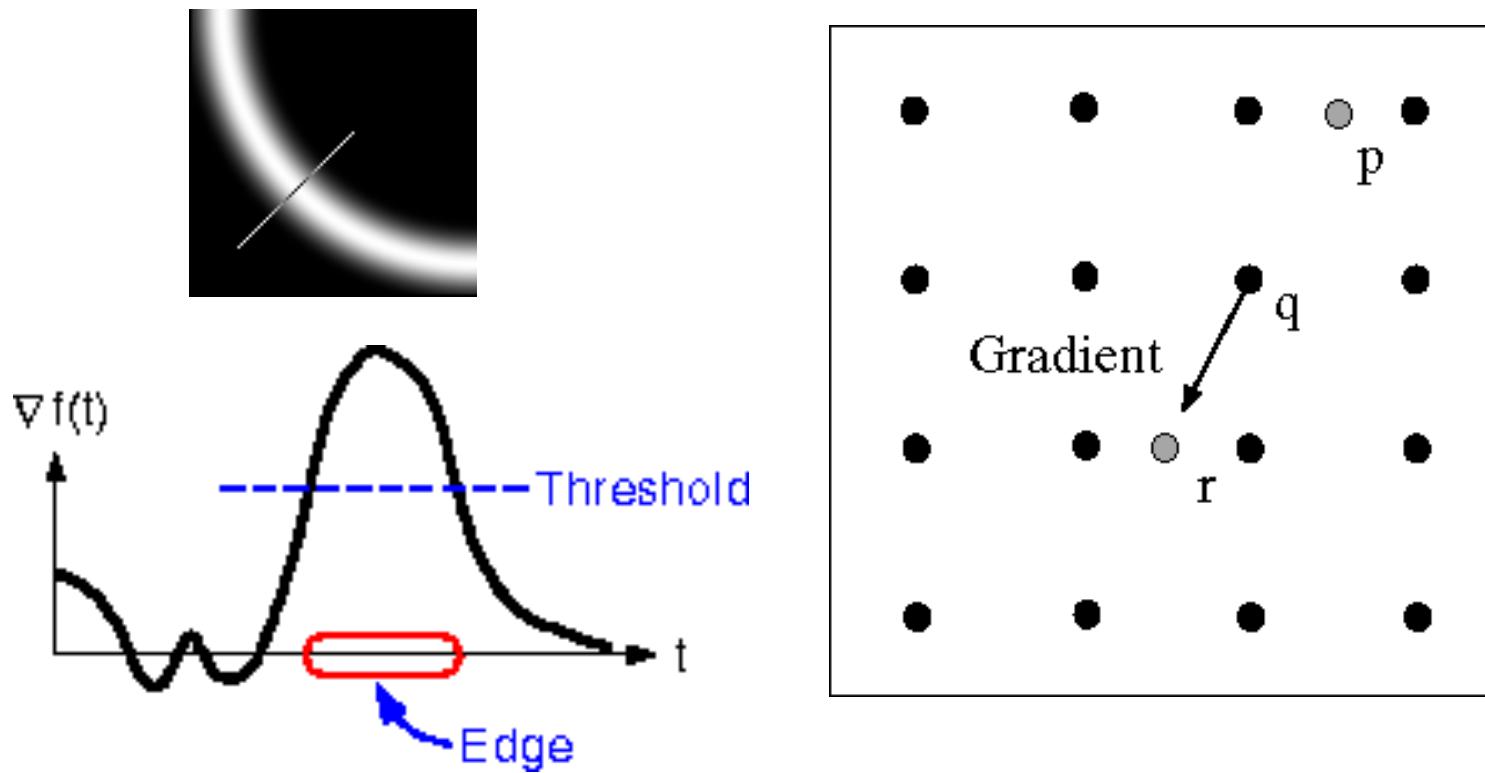
Building an edge detector



How to turn these thick regions of the gradient into curves?

Thresholded norm of the gradient

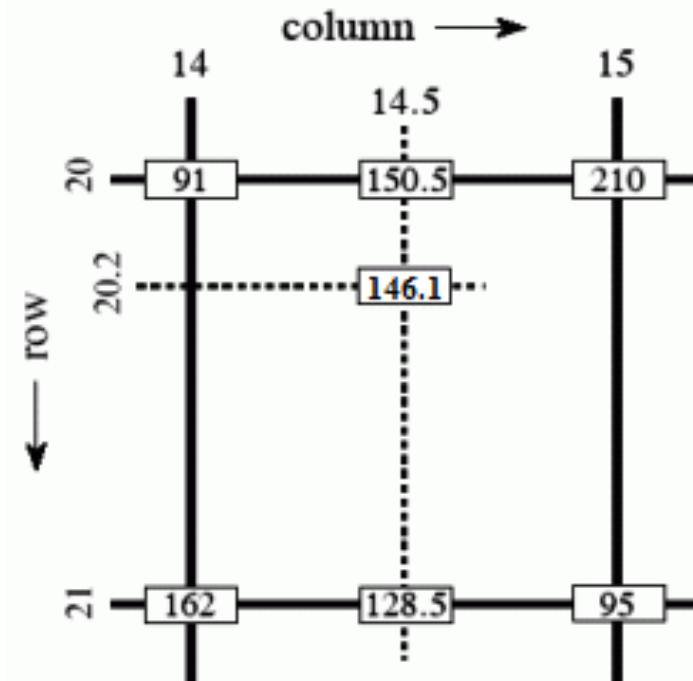
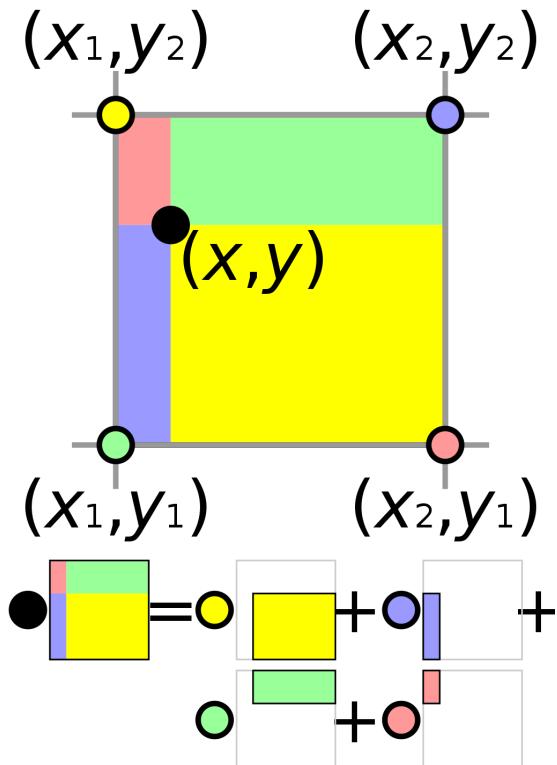
Non-maximum suppression



- For each location q above threshold, check that the gradient magnitude is higher than at neighbors p and r along the direction of the gradient
 - May need to interpolate to get the magnitudes at p and r

Bilinear Interpolation

$$f(x, y) \approx [1 - x \quad x] \begin{bmatrix} f(0,0) & f(0,1) \\ f(1,0) & f(1,1) \end{bmatrix} \begin{bmatrix} 1 - y \\ y \end{bmatrix}$$

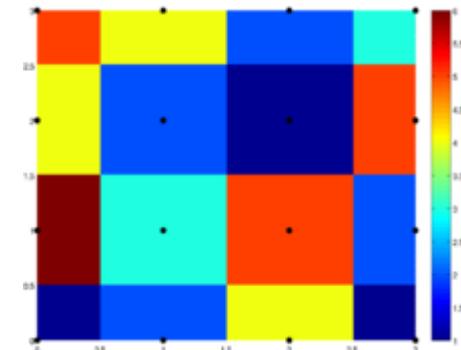


Sidebar: Interpolation options

```
imx2 = imresize(im, 2,  
    interpolation_type)
```

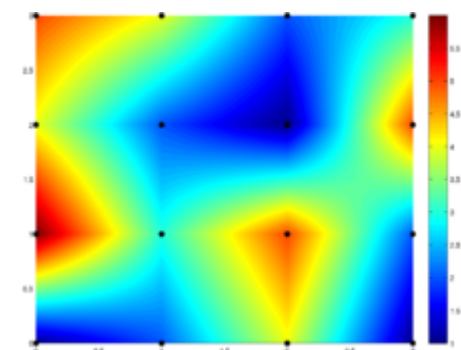
'nearest'

- Copy value from nearest known
- Very fast but creates blocky edges



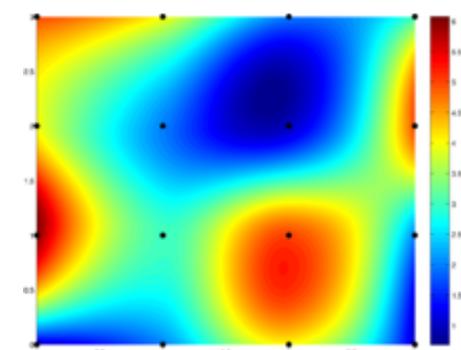
'bilinear'

- Weighted average from four nearest known pixels
- Fast and reasonable results



'bicubic' (default)

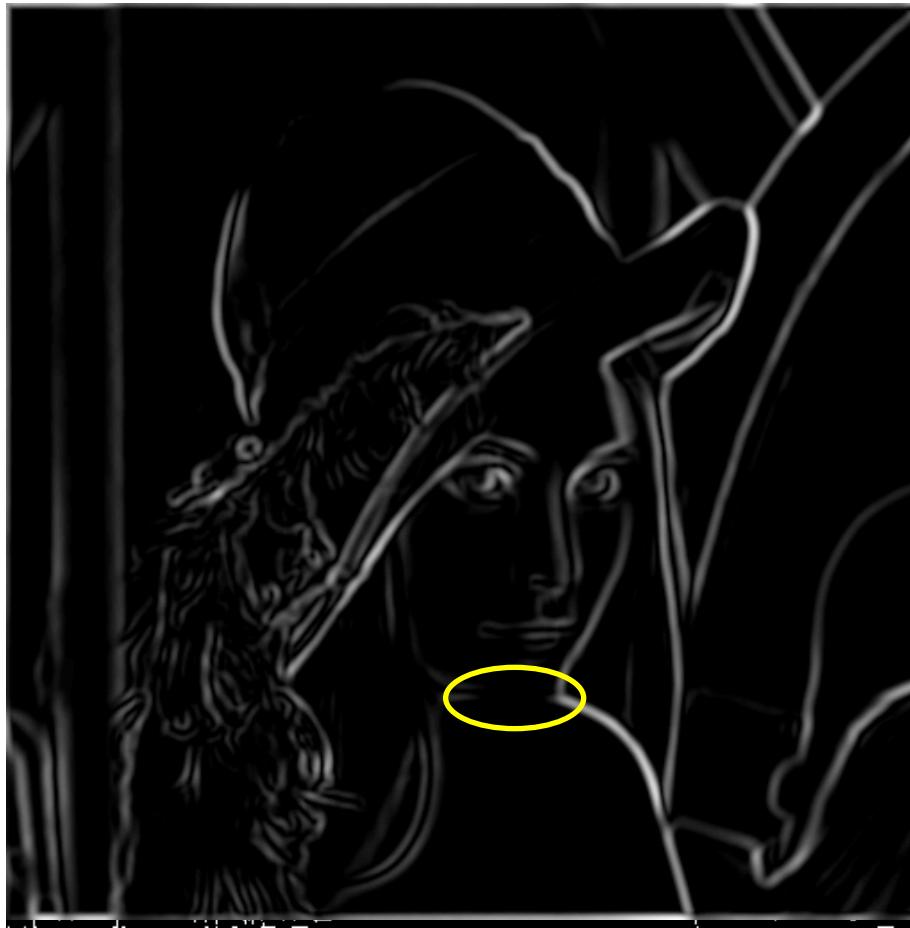
- Non-linear smoothing over larger area
- Slower, visually appealing, may create negative pixel values



Source: D. Hoeim

Examples from http://en.wikipedia.org/wiki/Bicubic_interpolation

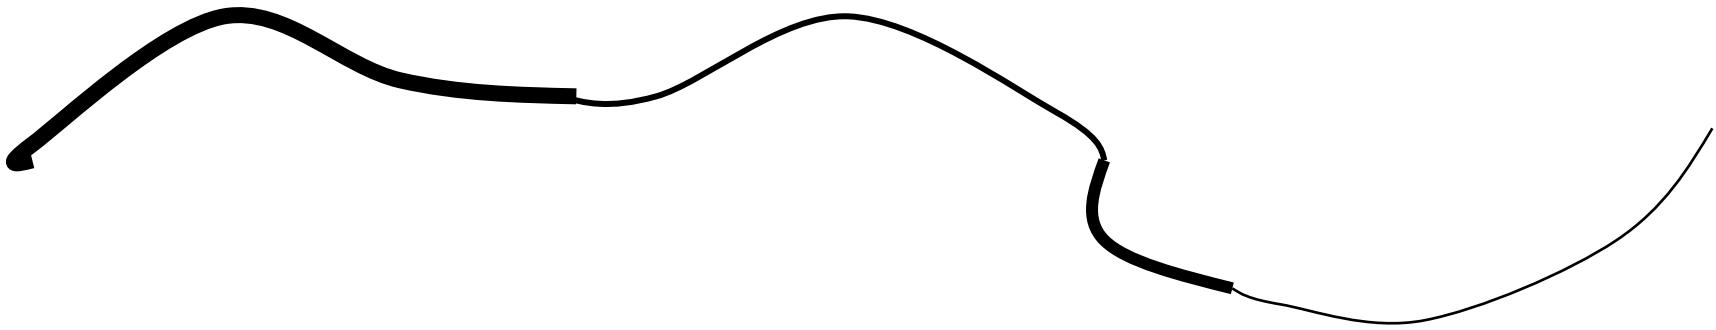
Non-maximum suppression



Another problem: pixels along this edge didn't survive the thresholding

Hysteresis thresholding

Use a high threshold to start edge curves, and a low threshold to continue them.



Hysteresis thresholding



original image



high threshold
(strong edges)



low threshold
(weak edges)



hysteresis threshold

Effect of σ (Gaussian kernel spread/size)



original



Canny with $\sigma = 1$



Canny with $\sigma = 2$

The choice of σ depends on desired behavior

- large σ detects large scale edges
- small σ detects fine features

Recap: Canny edge detector

1. Compute x and y gradient images
2. Find magnitude and orientation of gradient
3. **Non-maximum suppression:**
 - Thin wide “ridges” down to single pixel width
4. **Linking and thresholding (hysteresis):**
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them

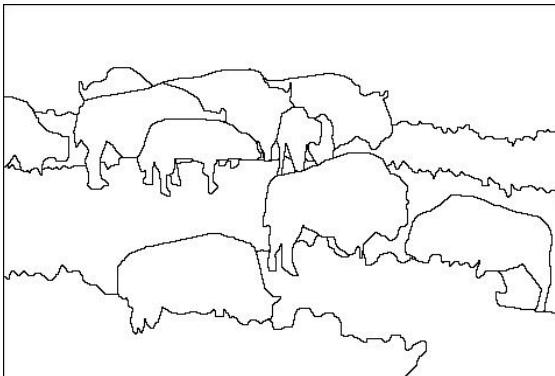
J. Canny, [A Computational Approach To Edge Detection](#), IEEE Trans. Pattern Analysis and Machine Intelligence, 8:679-714, 1986.

Image gradients vs. meaningful contours

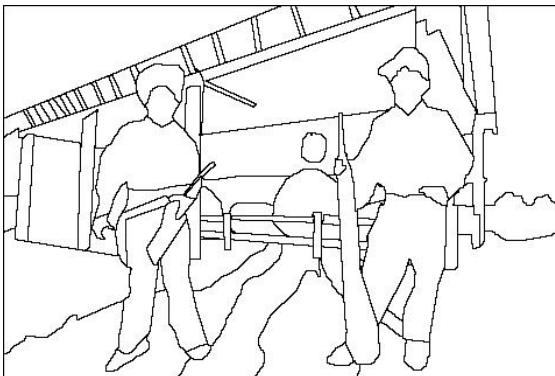
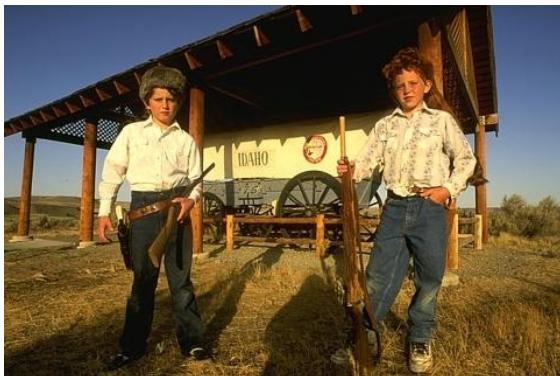
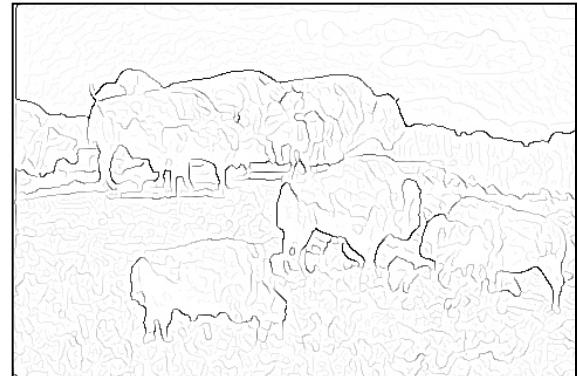
image



human segmentation



gradient magnitude

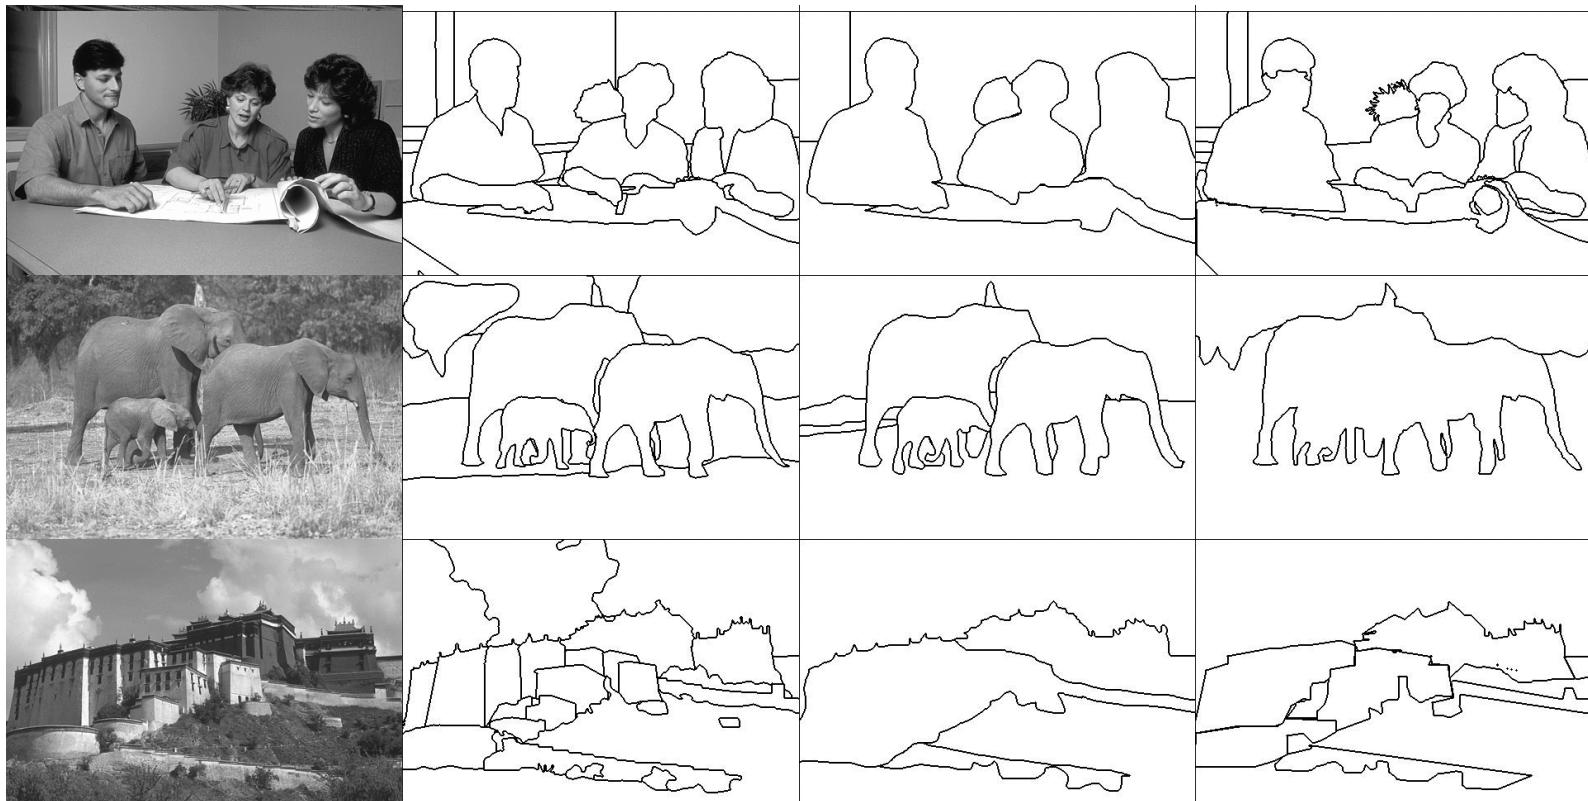


Berkeley segmentation database:

<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/segbench/>

Do humans consistently segment images?

Divide each image into pieces, where each piece represents a distinguished thing in the image. It is important that all of the pieces have approximately equal importance. The number of things in each image is up to you. Something between 2 and 20 should be reasonable for any of our images.



[A Database of Human Segmented Natural Images and its Application to Evaluating Segmentation Algorithms and Measuring Ecological Statistics](#)

Aside: Datasets, Metrics and Benchmarks

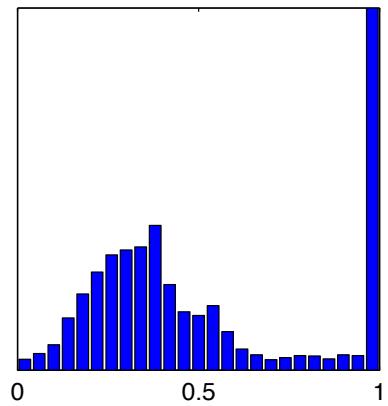
- Standard image sets
- Standard metrics
- Possible to *quantitatively* compare different methods

[A Database of Human Segmented Natural Images and its Application to Evaluating Segmentation Algorithms and Measuring Ecological Statistics](#)

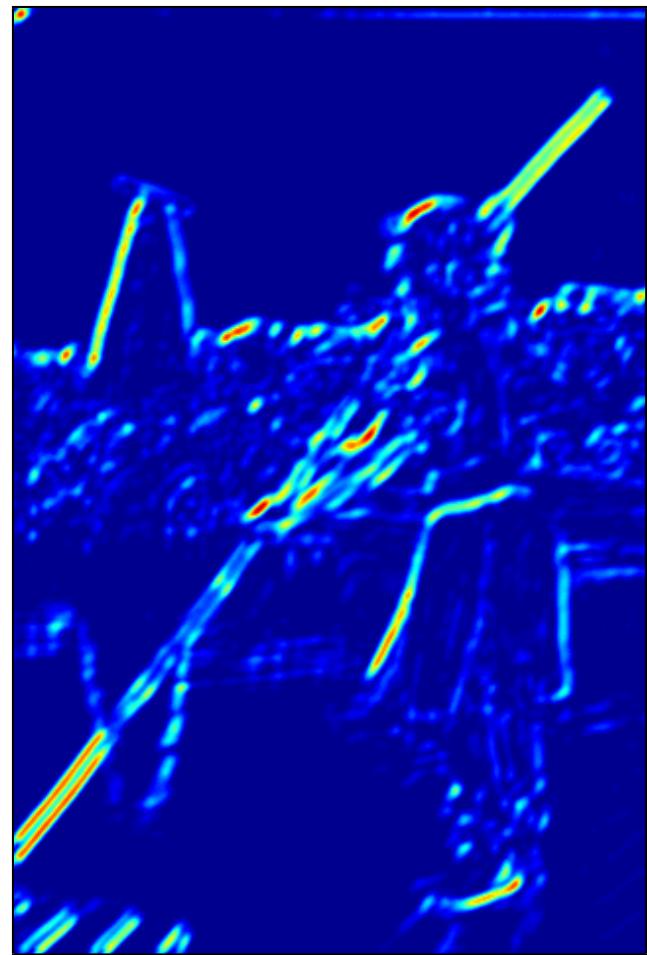
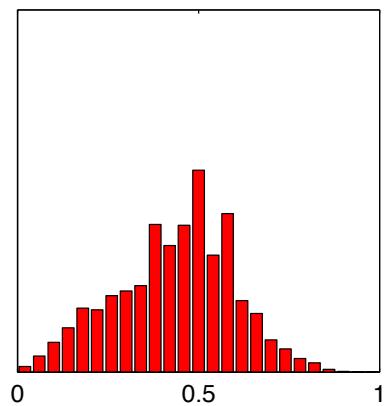
pB Boundary Detector

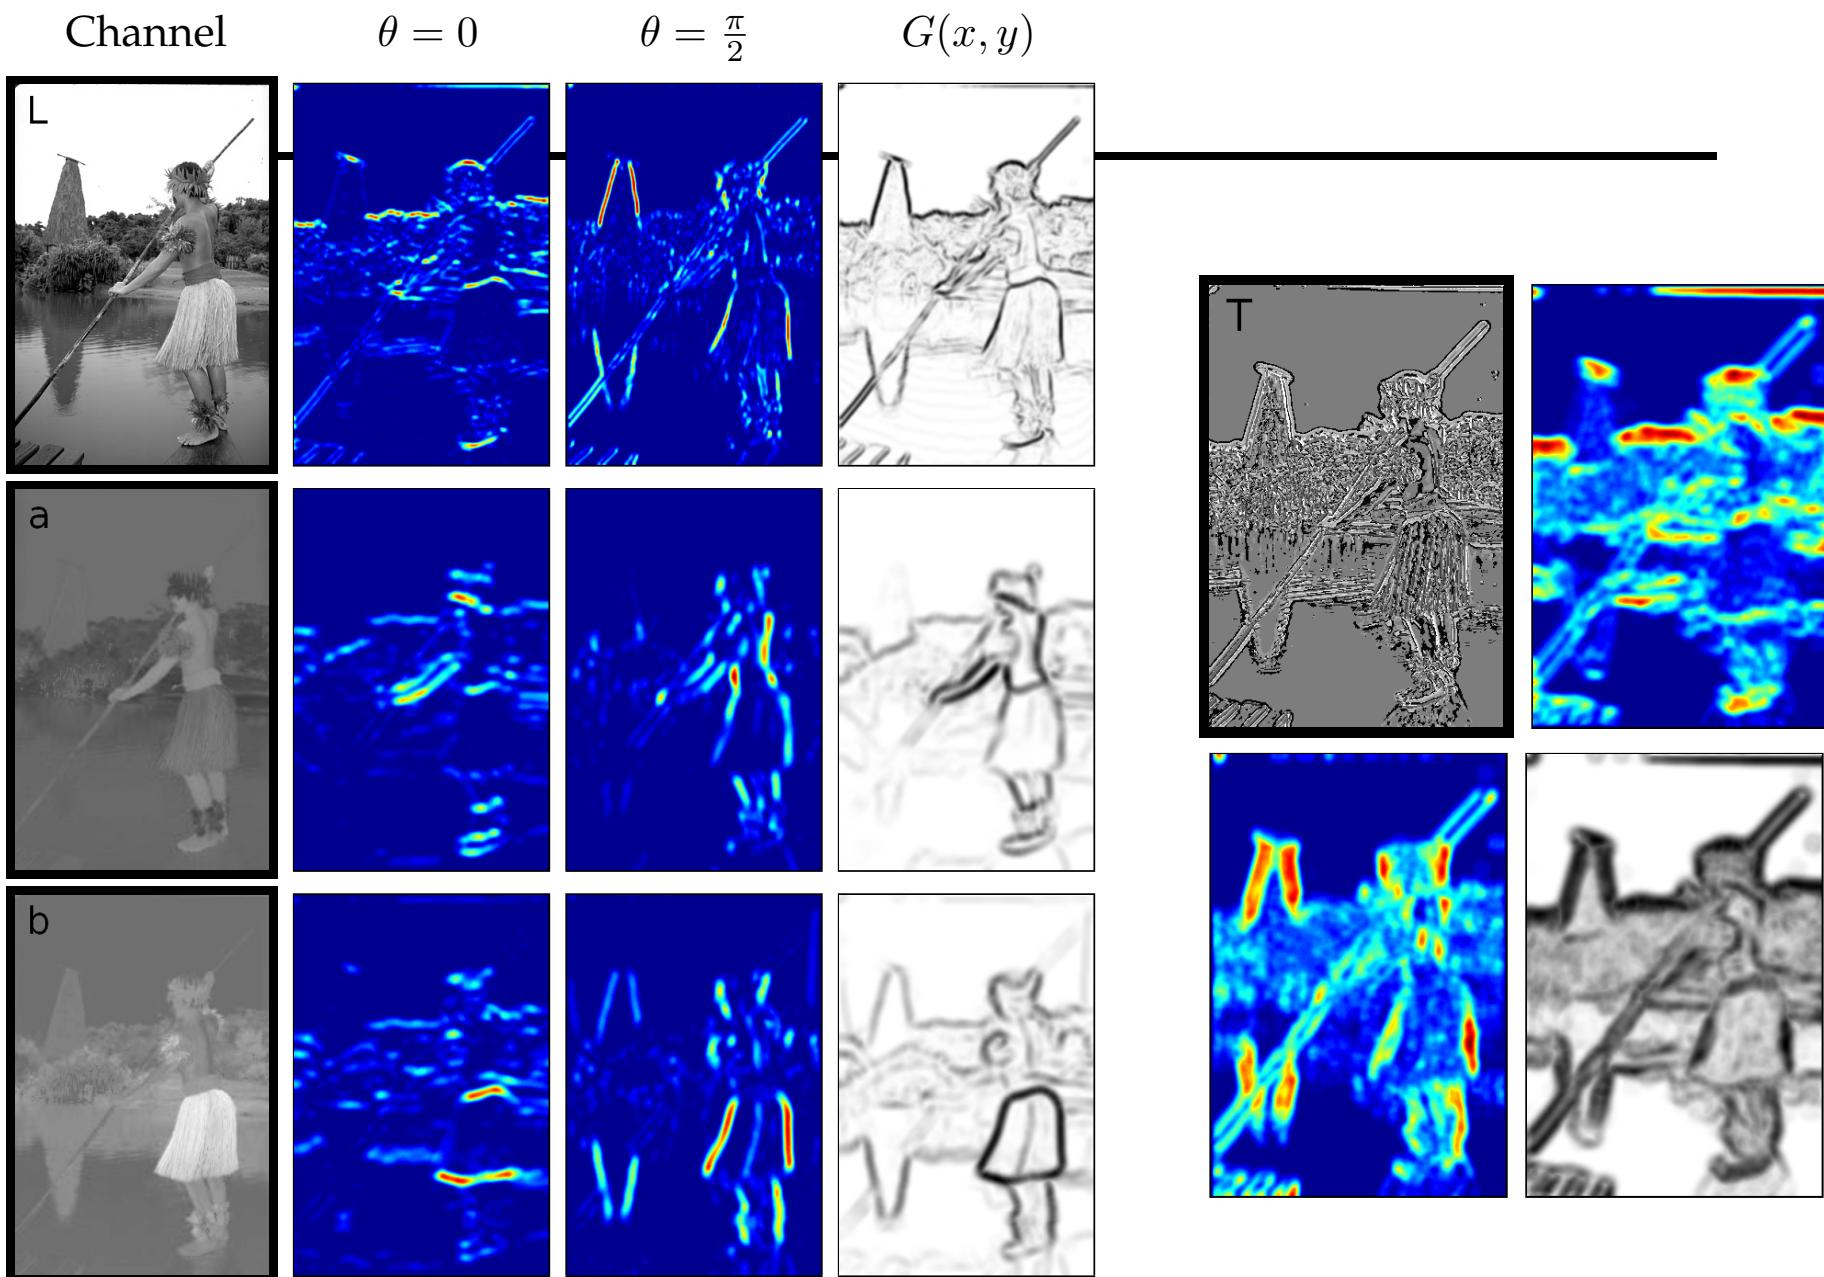


Upper Half-Disc Histogram



Lower Half-Disc Histogram





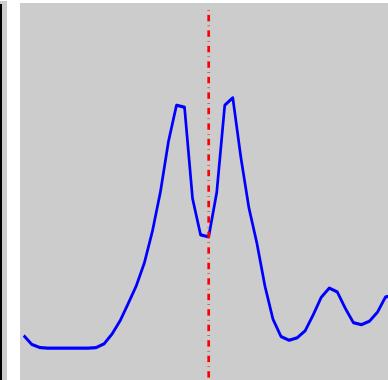
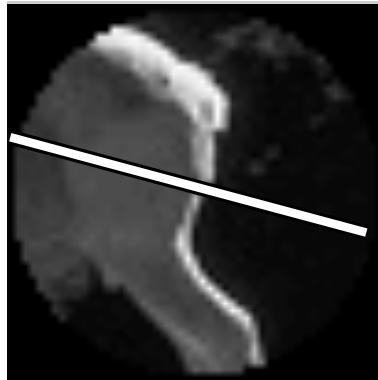
Lots of Tricks



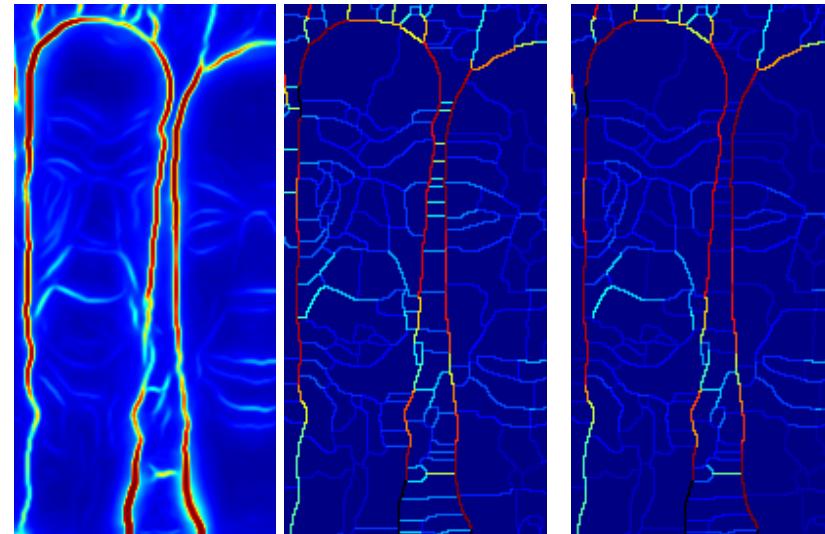
Sky Breaking



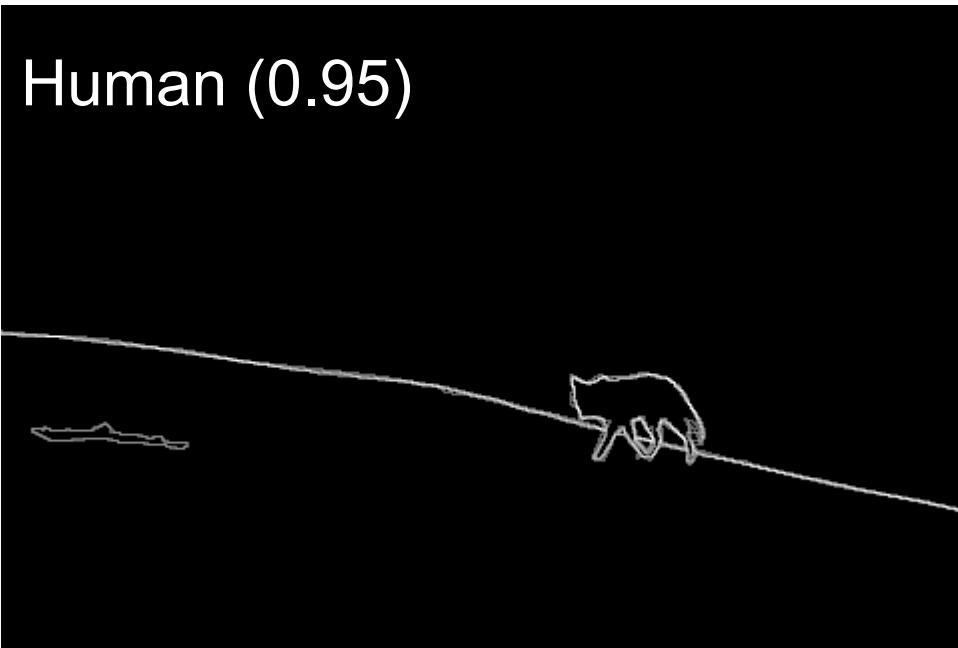
Information Leakage across Orientations



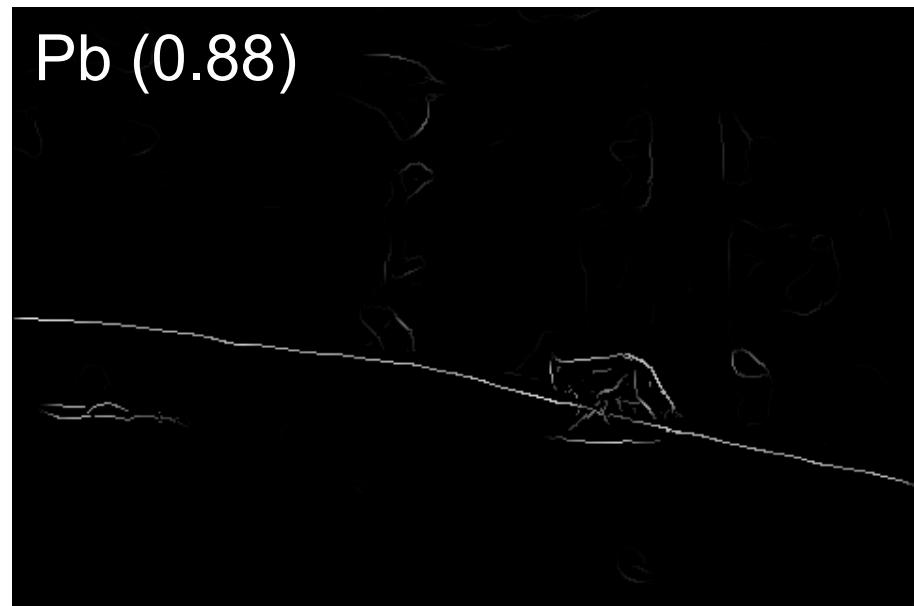
“Bird Edge” Problem with Texture Gradients



Results



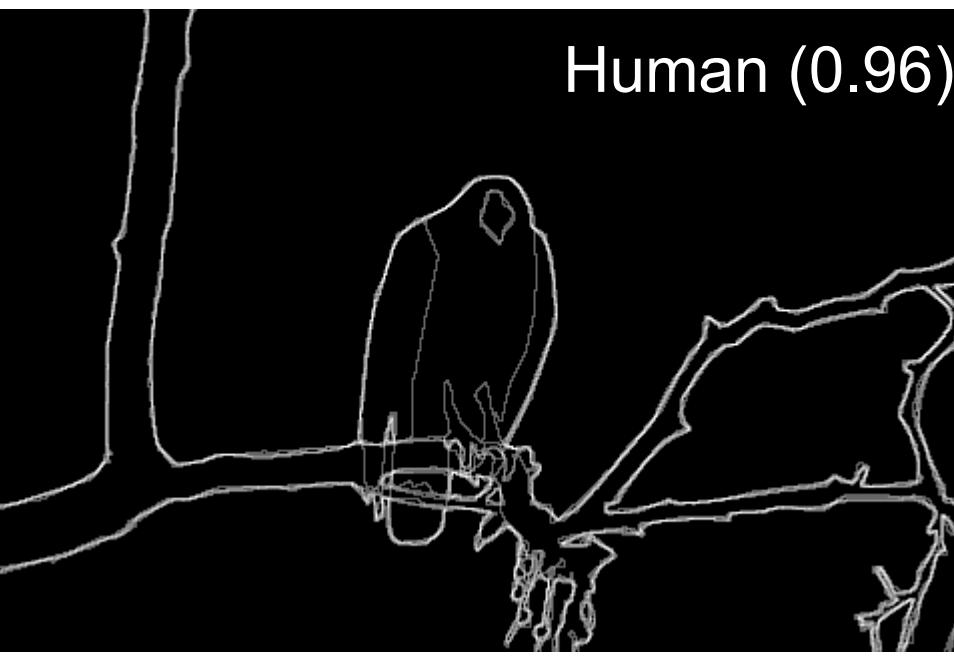
Pb (0.88)



Results



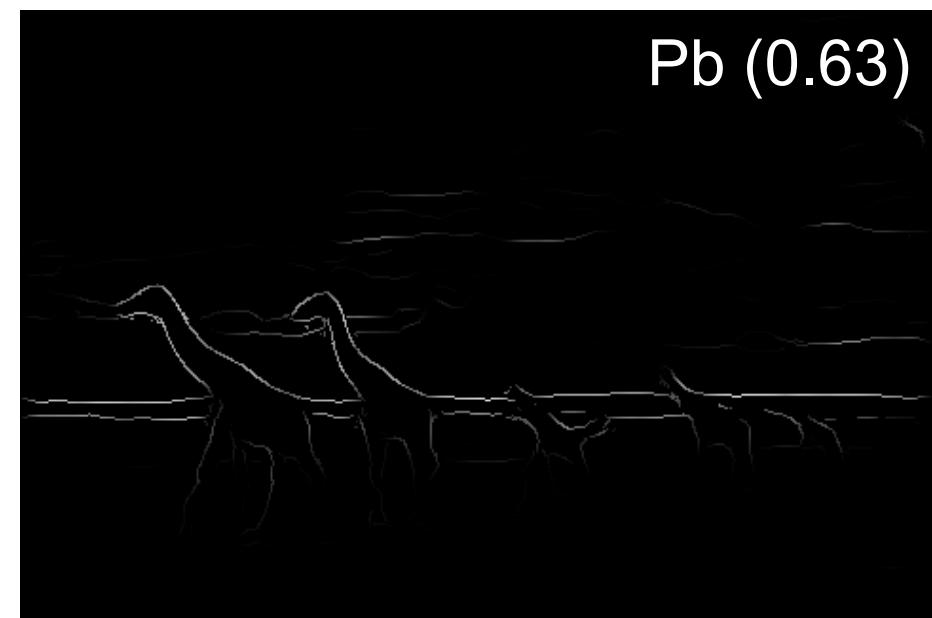
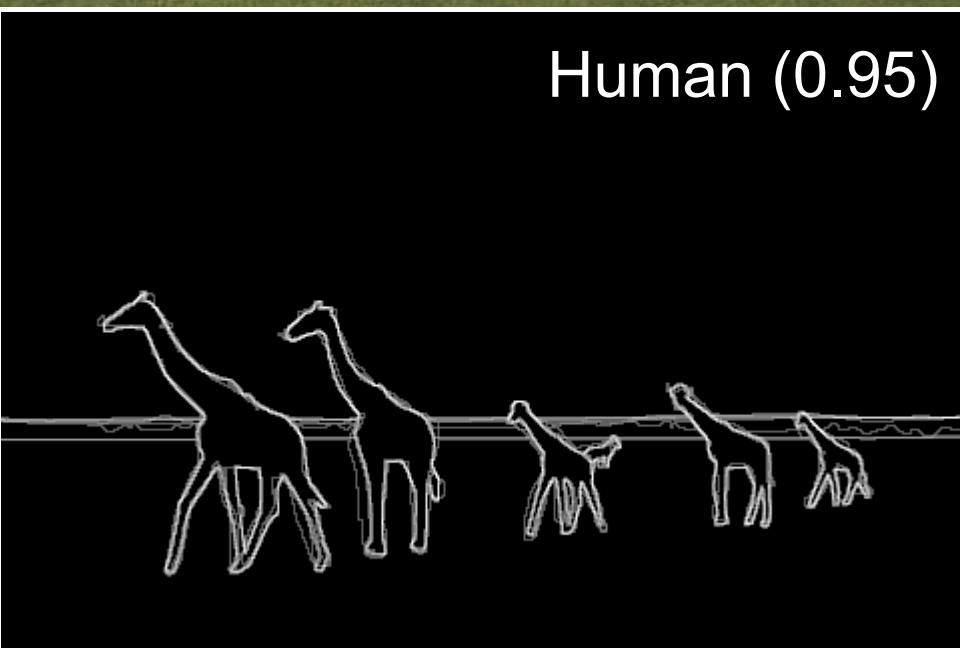
Human (0.96)



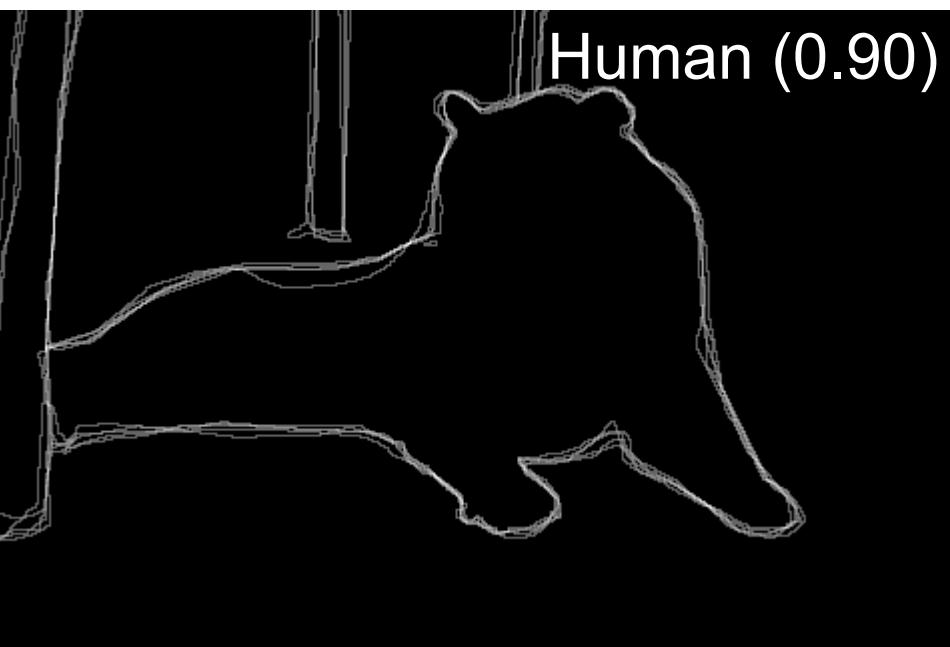
Pb (0.88)



Human (0.95)



Pb (0.63)

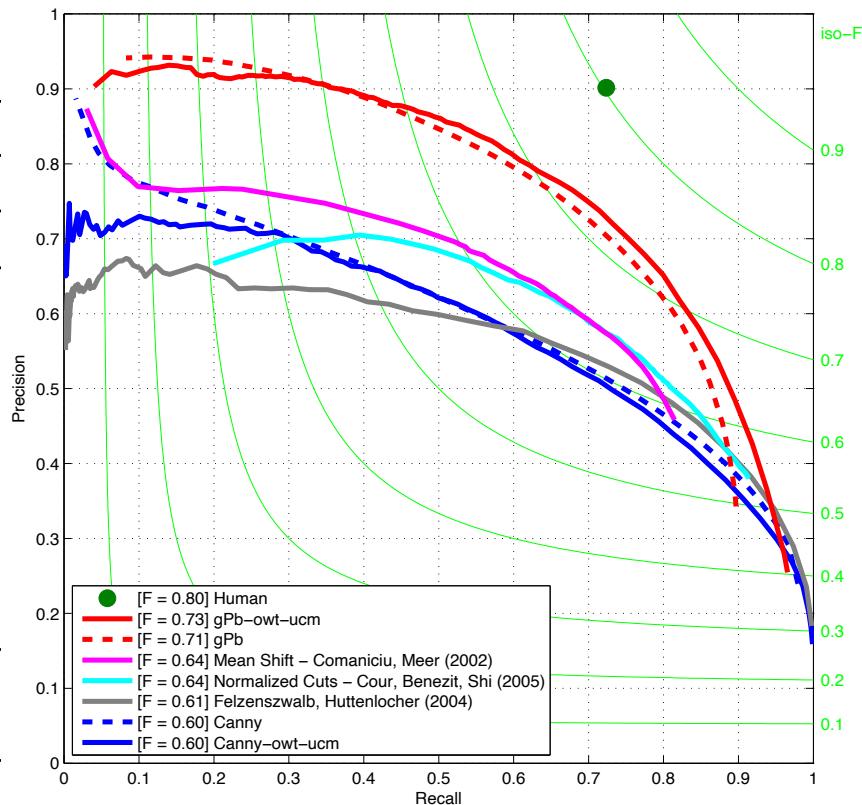


For more:

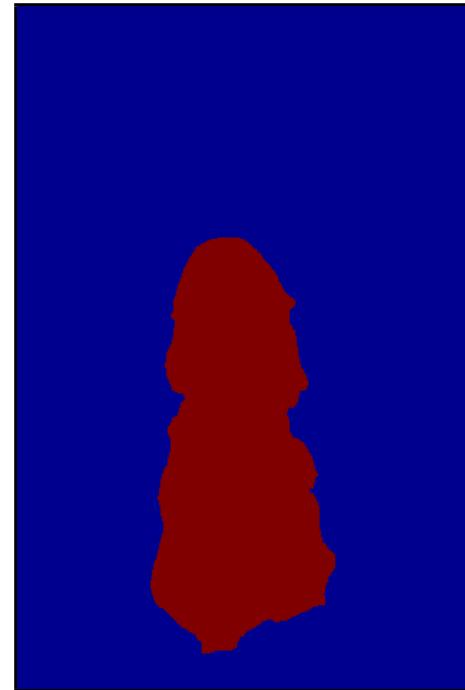
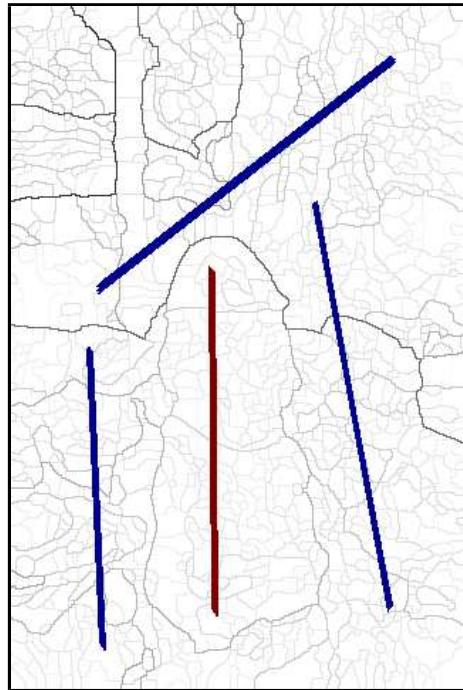
<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/bsds/bench/html/108082-color.html>

Empirical Research

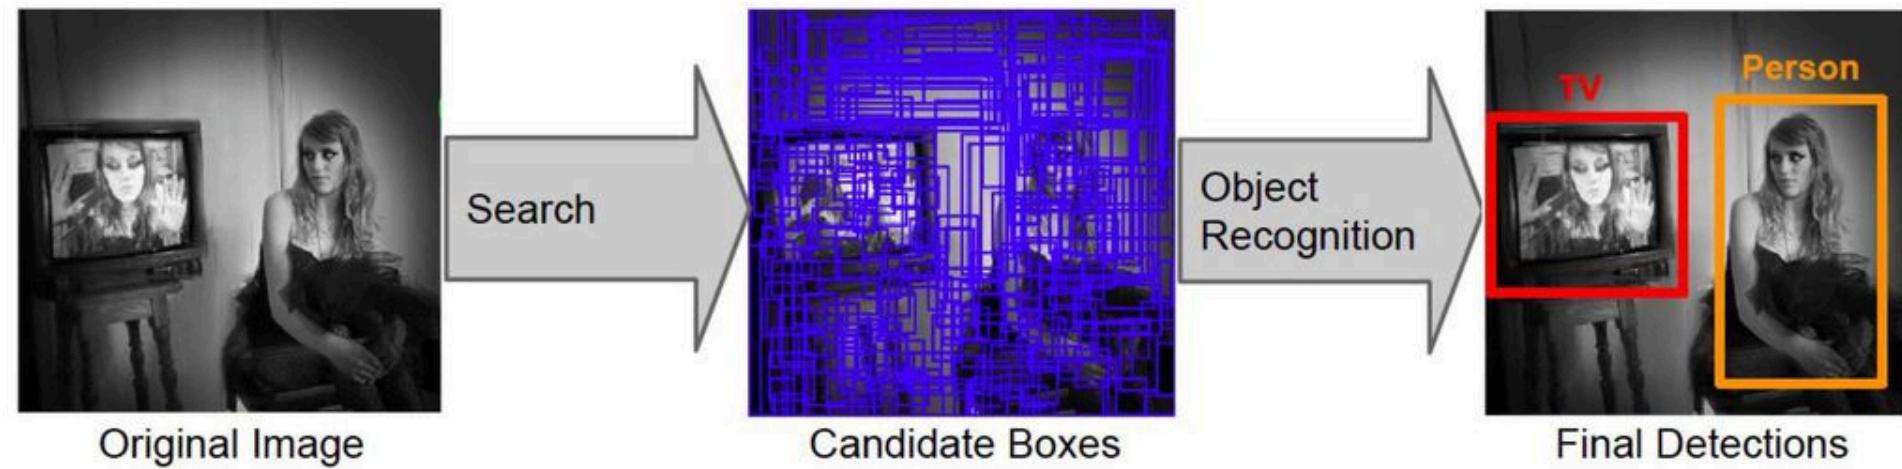
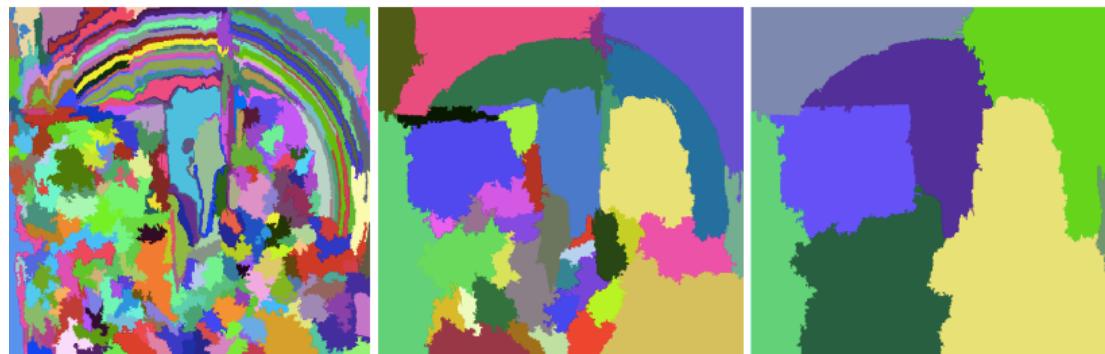
BSDS300			
	ODS	OIS	AP
Human	0.79	0.79	—
gPb-owt-ucm	0.71	0.74	0.73
[34] Mean Shift	0.63	0.66	0.54
[33] NCuts	0.62	0.66	0.43
Canny-owt-ucm	0.58	0.63	0.58
[32] Felz-Hutt	0.58	0.62	0.53
[31] SWA	0.56	0.59	0.54
Quad-Tree	0.37	0.39	0.26
gPb	0.70	0.72	0.66
Canny	0.58	0.62	0.58



Applications: Interactive Segmentation

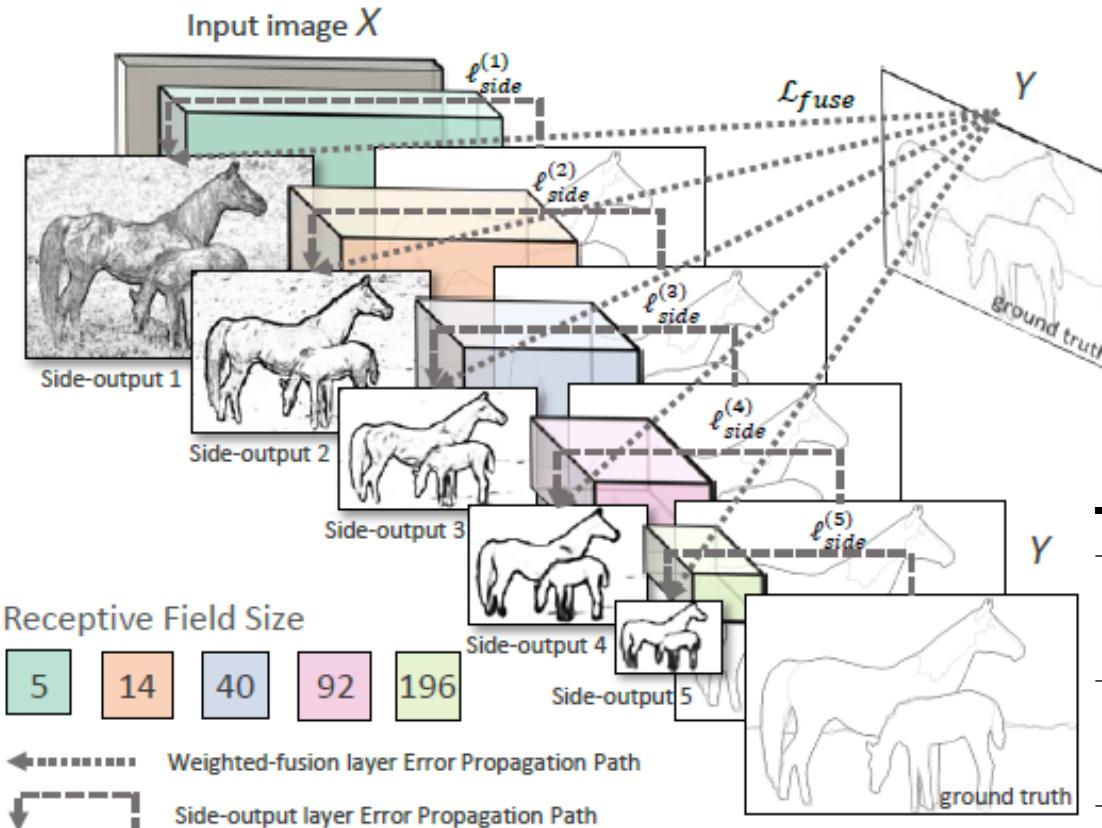


Applications: Pre-processing for Object Detection



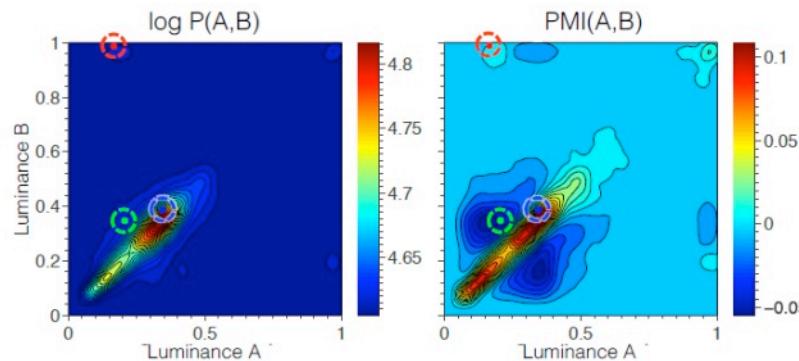
J. Uijlings, K. van de Sande, T. Gevers, and A. Smeulders,
[Selective Search for Object Recognition](#), IJCV 2013

Holistically nested edge detection



	ODS	OIS	AP	FPS
Human	.80	.80	-	-
Canny	.600	.640	.580	15
Felz-Hutt [9]	.610	.640	.560	10
BEL [5]	.660*	-	-	1/10
gPb-owt-ucm [1]	.726	.757	.696	1/240
Sketch Tokens [24]	.727	.746	.780	1
SCG [31]	.739	.758	.773	1/280
SE-Var [6]	.746	.767	.803	2.5
OEF [13]	.749	.772	.817	-
DeepNets [21]	.738	.759	.758	1/5†
N4-Fields [10]	.753	.769	.784	1/6†
DeepEdge [2]	.753	.772	.807	1/10 ³ †
CSCNN [19]	.756	.775	.798	-
DeepContour [34]	.756	.773	.797	1/30†
HED (ours)	.782	.804	.833	2.5†, 1/12

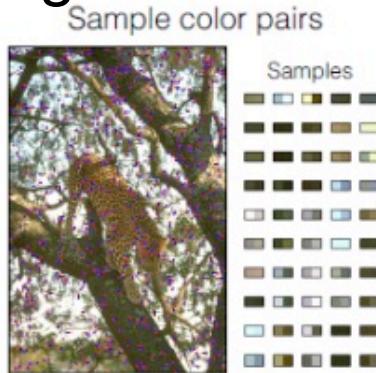
Crisp Boundary Detection using Pointwise Mutual Information (Isola et al. ECCV 2014)



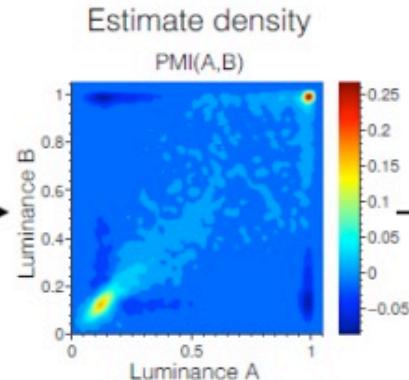
$$\text{PMI}_\rho(A, B) = \log \frac{P(A, B)^\rho}{P(A)P(B)}$$

Pixel combinations that are unlikely to be together are edges

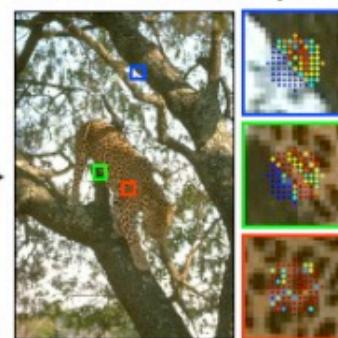
Algorithm:



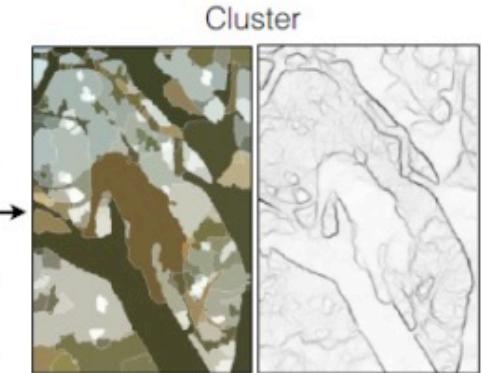
Kernel density estimation



Measure affinity

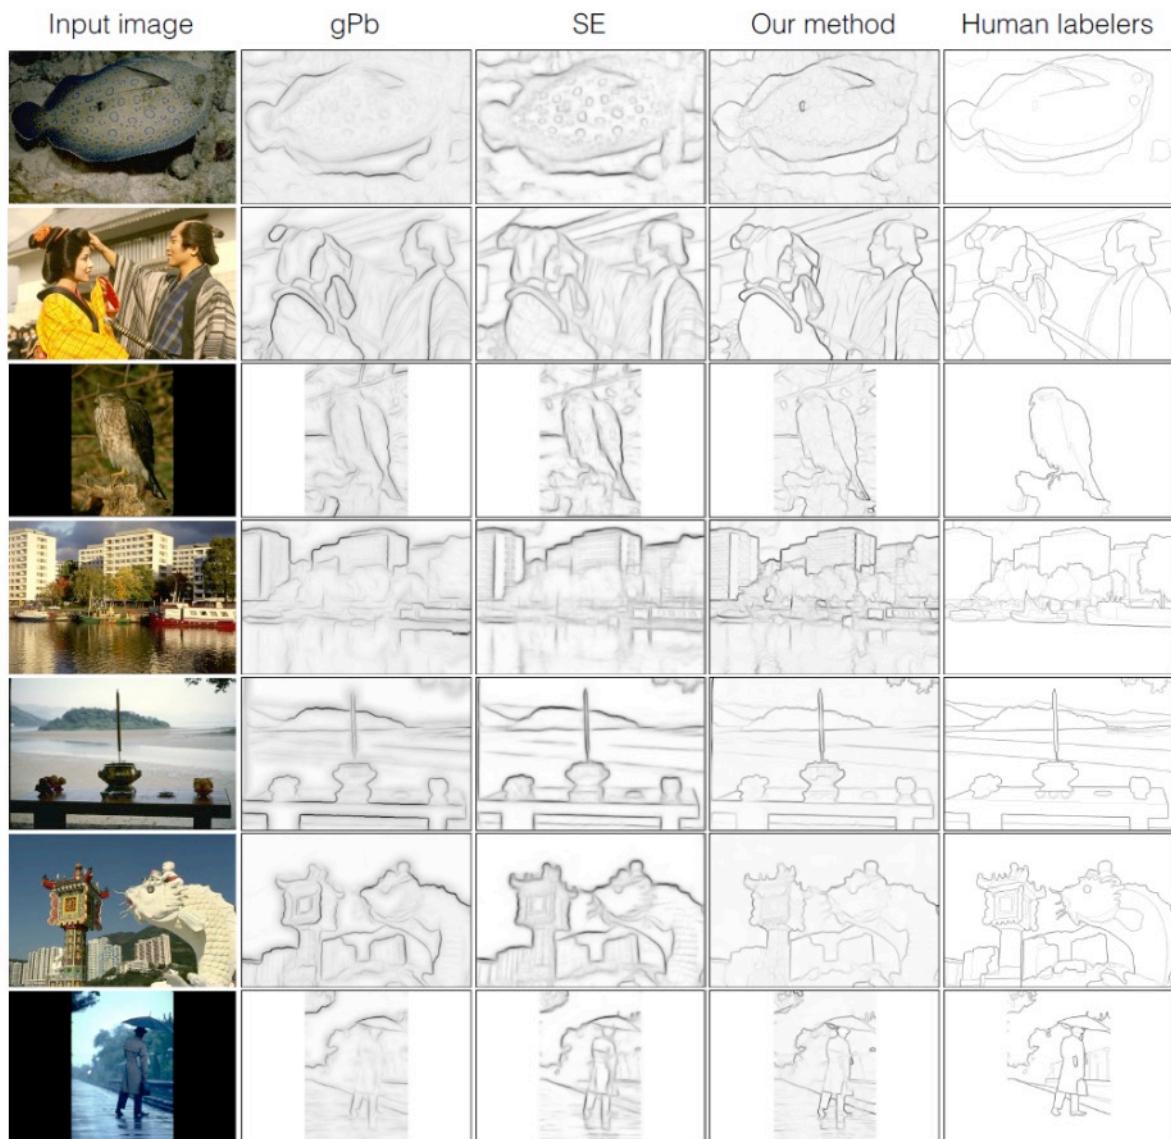


Spectral clustering



Crisp Boundary Detection using Pointwise Mutual Information

Algorithm	ODS	OIS	AP
Canny [14]	0.60	0.63	0.58
Mean Shift [36]	0.64	0.68	0.56
NCuts [37]	0.64	0.68	0.45
Felz-Hutt [38]	0.61	0.64	0.56
gPb [1]	0.71	0.74	0.65
gPb-owt-ucm [1]	0.73	0.76	0.73
SCG [9]	0.74	0.76	0.77
Sketch Tokens [7]	0.73	0.75	0.78
SE [8]	0.74	0.76	0.78
Our method – SS, color only	0.72	0.75	0.77
Our method – SS	0.73	0.76	0.79
Our method – MS	0.74	0.77	0.78



State of edge detection

Local edge detection is mostly solved

- Intensity gradient, color, texture
- HED on BSDS 500 is near human performance

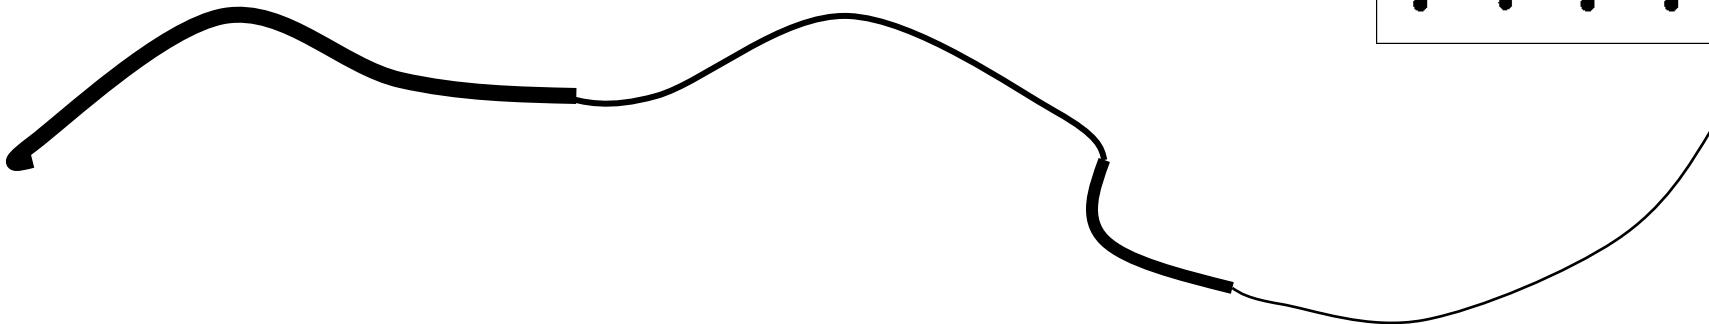
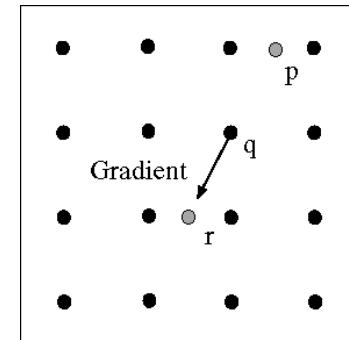
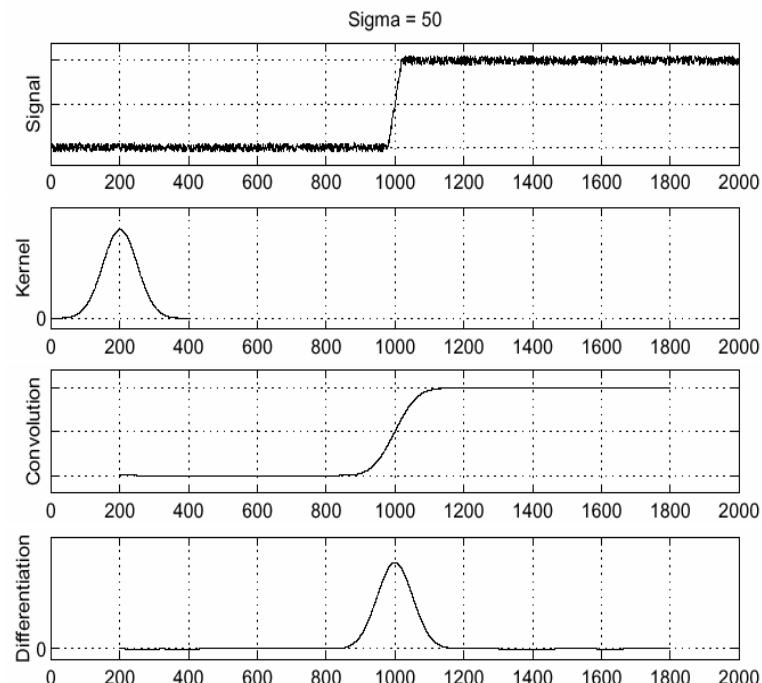
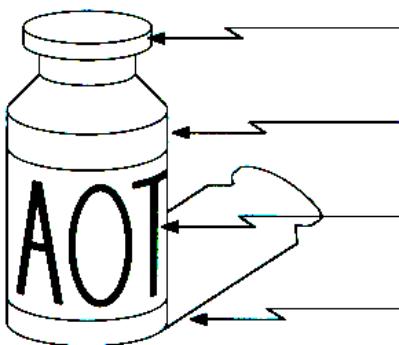
Some room for improvement by taking advantage of higher-level knowledge (e.g., objects)

Still hard to produce all objects within a small number of regions

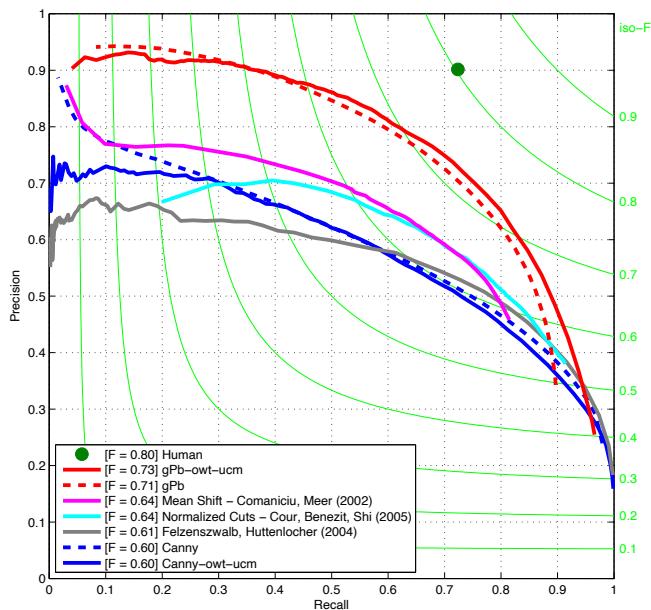
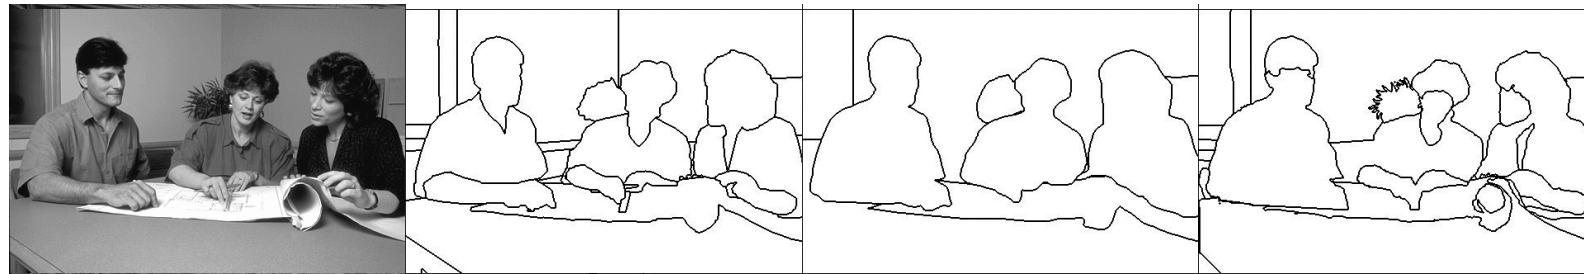
Are Edges an Input or an Output?



Recap



Recap



BSDS300			
	ODS	OIS	AP
Human	0.79	0.79	—
gPb-owt-ucm	0.71	0.74	0.73
[34] Mean Shift	0.63	0.66	0.54
[33] NCuts	0.62	0.66	0.43
Canny-owt-ucm	0.58	0.63	0.58
[32] Felz-Hutt	0.58	0.62	0.53
[31] SWA	0.56	0.59	0.54
Quad-Tree	0.37	0.39	0.26
gPb	0.70	0.72	0.66
Canny	0.58	0.62	0.58